

# Performance Analysis Report

## Single Source Shortest Path (SSSP) algorithms

Fakhir Ali i220762 Ayna Sulaiman i22105

### 1. Objective

The project aims to analyze the performance of **Single Source Shortest Path (SSSP)** algorithms under dynamic network updates. Implementations were tested across three environments:

- **Serial (Baseline)**
- **OpenMP (Shared Memory Parallelism)**
- **MPI + OpenMP (Distributed Memory with Intra-node Parallelism)**

The goal is to evaluate the efficiency of dynamic update strategies (insertions, deletions, or both) against full recomputation, in terms of runtime and scalability.

### 2. Experimental Setup

- **Datasets:** `amazon400.txt`, `verybig.txt`
- **Platforms:**
  - OpenMP: Multi-threaded shared-memory CPU
  - MPI: 2-process distributed execution using `part0.txt`, `part1.txt`
- **Update Patterns:** Varying combinations of edge insertions and deletions
- **Metrics:**
  - Initial SSSP time
  - Recompute time after updates
  - Dynamic async update time
  - Async depth for update propagation control

### 3. Serial Implementation Analysis

#### Key Findings

Dataset	Insertions	Deletions	Recompute Time (ms)	Async Update Time (ms)	Speedup
amazon400.txt	100,000	0	640	550	1.16×
amazon400.txt	10,000	0	580	170	3.41×
amazon400.txt	0	10,000	566	104	5.44×
verybig.txt	1,000,000	50,000	13,991	4,507	3.10×
verybig.txt	50,000	1,000,000	13,443	2,405	5.59×

#### Observations

- Serial dynamic updates are faster than recomputation for small- to medium-scale changes.
- Performance gain decreases with higher insertion counts or deep structural changes.
- Deletion-heavy updates impact the tree more but still benefit from selective recomputation.

### 4. OpenMP Implementation Analysis

#### Key Findings

Scenario	Insertions	Deletions	Async Depth	Recompute Time (ms)	Async Update Time (ms)	Speedup
A	1,000,000	50,000	5	12,899	1,899	6.79×
B	1,000,000	50,000	5	13,574	346	39.2×
C	50,000	1,000,000	1	13,502	1,913	7.06×
D	50,000	1,000,000	1	11,572	330	35.0×

#### Observations

- OpenMP provides a large speedup (5–39×) over recomputation, especially for insert-heavy updates.
- Async depth influences performance. Deeper depths allow for broader update propagation and faster convergence.
- Batched processing of updates and depth-bounded parallel traversals improve load balancing and reduce overhead

## 5. MPI + OpenMP Hybrid Implementation Analysis

### Execution Flow Summary

- MPI partitions the graph; each process handles a subset.
- Within each MPI process, OpenMP handles parallel SSSP update.
- `broadcast_vector()` syncs initial SSSP results across processes.
- `exchange_ghost_distances()` ensures consistency of boundary node distances.
- Dynamic updates (insertion/deletion) are processed locally and asynchronously in parallel.

### Expected Behavior

Factor	Contribution to Performance
Inter-node communication	Incurred for boundary node updates
Intra-node OpenMP	Enhances local traversal efficiency
Async update strategy	Limits redundant work per iteration
Ghost sync + AllReduce	Ensures convergence across partitions

### Performance Insights

While specific MPI timings were not provided, the hybrid code is designed to:

- Minimize communication using boundary sync.
- Use OpenMP for parallel update propagation.
- Handle bulk updates efficiently in distributed memory.
- Avoid global locks or barriers (only uses `MPI_Allreduce`).

The hybrid approach is expected to perform well for large-scale graphs and is scalable to multiple processes.

## 6. Comparative Summary

Method	Scalability	Best Use Case	Speedup Over Recompute	Notes
Serial	Limited	Small-scale graphs, few updates	Up to 5.6×	Simple, easy to debug
OpenMP	Moderate–High	Shared memory systems	Up to 39×	Fastest async update among all
MPI+OpenMP	High (distributed)	Large-scale graphs, distributed nodes	Not benchmarked here	Supports inter-node parallelism & scaling

## 7. Conclusions

- **Dynamic updates are consistently faster than full recomputation** across all implementations, especially with OpenMP.
- **OpenMP** offers excellent speedups and is suitable for shared-memory environments with large graphs.
- The **hybrid MPI+OpenMP** model supports scalability and efficient distributed processing and is well-aligned with the paper’s objectives.
- Results are consistent with the referenced research paper, confirming both the validity and performance advantage of the parallel updating strategy over static recomputation.